

```
pip install openpyxl
```

Requirement already satisfied: openpyxl in c:\users\gaikw\appdata\local\programs\python\python311\lib\site-packages (3.1.5)

Requirement already satisfied: et-xmlfile in c:\users\gaikw\appdata\local\programs\python\python311\lib\site-packages (from openpyxl) (2.0.0)

Note: you may need to restart the kernel to use updated packages.

[notice] A new release of pip available: 22.3 -> 24.3.1

[notice] To update, run: python.exe -m pip install --upgrade pip

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
file_path = "ENB2012_data.xlsx"
data = pd.read_excel(file_path)
```

```
df = pd.DataFrame(data)
df.info()
df.isna().sum()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 10 columns):
```

| # | Column | Non-Null Count | Dtype |
|---|--------|----------------|---------|
| 0 | X1 | 768 non-null | float64 |
| 1 | X2 | 768 non-null | float64 |
| 2 | X3 | 768 non-null | float64 |
| 3 | X4 | 768 non-null | float64 |
| 4 | X5 | 768 non-null | float64 |
| 5 | X6 | 768 non-null | int64 |
| 6 | X7 | 768 non-null | float64 |
| 7 | X8 | 768 non-null | int64 |
| 8 | Y1 | 768 non-null | float64 |
| 9 | Y2 | 768 non-null | float64 |

```
dtypes: float64(8), int64(2)
```

```
memory usage: 60.1 KB
```

```
X1    0
X2    0
X3    0
X4    0
X5    0
X6    0
```

```
X7    0
X8    0
Y1    0
Y2    0
dtype: int64
```

```
df2 = df.dropna()
print(df2)
df2.isnull().sum()
```

| | X1 | X2 | X3 | X4 | X5 | X6 | X7 | X8 | Y1 | Y2 |
|-----|------|-------|-------|--------|-----|-----|-----|-----|-------|-------|
| 0 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | 2 | 0.0 | 0 | 15.55 | 21.33 |
| 1 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | 3 | 0.0 | 0 | 15.55 | 21.33 |
| 2 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | 4 | 0.0 | 0 | 15.55 | 21.33 |
| 3 | 0.98 | 514.5 | 294.0 | 110.25 | 7.0 | 5 | 0.0 | 0 | 15.55 | 21.33 |
| 4 | 0.90 | 563.5 | 318.5 | 122.50 | 7.0 | 2 | 0.0 | 0 | 20.84 | 28.28 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 0.64 | 784.0 | 343.0 | 220.50 | 3.5 | 5 | 0.4 | 5 | 17.88 | 21.40 |
| 764 | 0.62 | 808.5 | 367.5 | 220.50 | 3.5 | 2 | 0.4 | 5 | 16.54 | 16.88 |
| 765 | 0.62 | 808.5 | 367.5 | 220.50 | 3.5 | 3 | 0.4 | 5 | 16.44 | 17.11 |
| 766 | 0.62 | 808.5 | 367.5 | 220.50 | 3.5 | 4 | 0.4 | 5 | 16.48 | 16.61 |
| 767 | 0.62 | 808.5 | 367.5 | 220.50 | 3.5 | 5 | 0.4 | 5 | 16.64 | 16.03 |

```
[768 rows x 10 columns]
```

```
X1    0
X2    0
X3    0
X4    0
X5    0
X6    0
X7    0
X8    0
Y1    0
Y2    0
dtype: int64
```

```
df2[df2.columns].corr()
```

| | X1 | X2 | X3 | X4 |
|------|---------------|---------------|---------------|---------------|
| X5 \ | | | | |
| X1 | 1.000000e+00 | -9.919015e-01 | -2.037817e-01 | -8.688234e-01 |
| X2 | -9.919015e-01 | 1.000000e+00 | 1.955016e-01 | 8.807195e-01 |
| X3 | -2.037817e-01 | 1.955016e-01 | 1.000000e+00 | -2.923165e-01 |
| X4 | -8.688234e-01 | 8.807195e-01 | -2.923165e-01 | 1.000000e+00 |
| X5 | 8.277473e-01 | -8.581477e-01 | 2.809757e-01 | -9.725122e-01 |

```

1.000000e+00
X6  4.678592e-17 -3.459372e-17 -2.429499e-17 -5.830058e-17  4.492205e-
17
X7 -2.960552e-15  3.636925e-15 -8.567455e-17 -1.759011e-15  1.489134e-
17
X8 -7.107006e-16  2.438409e-15  2.067384e-16 -1.078071e-15 -2.920613e-
17
Y1  6.222719e-01 -6.581199e-01  4.556714e-01 -8.618281e-01  8.894305e-
01
Y2  6.343391e-01 -6.729989e-01  4.271170e-01 -8.625466e-01  8.957852e-
01

```


| | X6 | X7 | X8 | Y1 | Y2 |
|----|---------------|---------------|---------------|-----------|-----------|
| X1 | 4.678592e-17 | -2.960552e-15 | -7.107006e-16 | 0.622272 | 0.634339 |
| X2 | -3.459372e-17 | 3.636925e-15 | 2.438409e-15 | -0.658120 | -0.672999 |
| X3 | -2.429499e-17 | -8.567455e-17 | 2.067384e-16 | 0.455671 | 0.427117 |
| X4 | -5.830058e-17 | -1.759011e-15 | -1.078071e-15 | -0.861828 | -0.862547 |
| X5 | 4.492205e-17 | 1.489134e-17 | -2.920613e-17 | 0.889430 | 0.895785 |
| X6 | 1.000000e+00 | -9.406007e-16 | -2.549352e-16 | -0.002587 | 0.014290 |
| X7 | -9.406007e-16 | 1.000000e+00 | 2.129642e-01 | 0.269842 | 0.207505 |
| X8 | -2.549352e-16 | 2.129642e-01 | 1.000000e+00 | 0.087368 | 0.050525 |
| Y1 | -2.586763e-03 | 2.698417e-01 | 8.736846e-02 | 1.000000 | 0.975862 |
| Y2 | 1.428960e-02 | 2.075050e-01 | 5.052512e-02 | 0.975862 | 1.000000 |

```

plt.figure(figsize =(15,15))
sns.heatmap(data=df2[df2.columns].corr(), annot = True, linewidth = 2,
linecolor = 'white', cmap='Blues')

```

<Axes: >



```
from sklearn.model_selection import train_test_split
#Ignoring the features having weak/no correlation.
features_col = ['X1', 'X2', 'X4', 'X5', 'X7']
target_col = ['Y1']
x = df2[features_col]
y = df2[target_col]

x_train, x_test, y_train, y_test = train_test_split(x, y,
random_state=64, test_size=0.2)
print('X Train :\n', x_train.head())
```

```

print('\nY Train :\n', y_train.head())
print('\nX Test :\n', x_test.head())
print('\nY Test :\n', y_test.head() , '\n')

```

```

x_train = x_train.to_numpy()
y_train = y_train.to_numpy()
x_test = x_test.to_numpy()
y_test = y_test.to_numpy()

```

```

print('Shape of x_train :', x_train.shape)
print('Shape of y_train :', y_train.shape)
print('Shape of x_test :', x_test.shape)
print('Shape of y_test :', y_test.shape)

```

X Train :

| | X1 | X2 | X4 | X5 | X7 |
|-----|------|-------|-------|-----|------|
| 588 | 0.82 | 612.5 | 147.0 | 7.0 | 0.40 |
| 128 | 0.69 | 735.0 | 220.5 | 3.5 | 0.10 |
| 443 | 0.86 | 588.0 | 147.0 | 7.0 | 0.25 |
| 644 | 0.76 | 661.5 | 122.5 | 7.0 | 0.40 |
| 127 | 0.71 | 710.5 | 220.5 | 3.5 | 0.10 |

Y Train :

| | Y1 |
|-----|-------|
| 588 | 28.95 |
| 128 | 11.45 |
| 443 | 29.88 |
| 644 | 39.32 |
| 127 | 10.68 |

X Test :

| | X1 | X2 | X4 | X5 | X7 |
|-----|------|-------|-------|-----|------|
| 270 | 0.71 | 710.5 | 220.5 | 3.5 | 0.10 |
| 445 | 0.82 | 612.5 | 147.0 | 7.0 | 0.25 |
| 93 | 0.62 | 808.5 | 220.5 | 3.5 | 0.10 |
| 670 | 0.62 | 808.5 | 220.5 | 3.5 | 0.40 |
| 236 | 0.62 | 808.5 | 220.5 | 3.5 | 0.10 |

Y Test :

| | Y1 |
|-----|-------|
| 270 | 10.67 |
| 445 | 24.96 |
| 93 | 12.97 |
| 670 | 16.55 |
| 236 | 12.85 |

```

Shape of x_train : (614, 5)
Shape of y_train : (614, 1)

```

Shape of x_test : (154, 5)

Shape of y_test : (154, 1)

```
def multilinear_regression(x, y):  
    b0 = np.ones(x.shape[0])  
    new_x = np.concatenate((b0.reshape(-1,1), x), axis = 1)  
    beta = np.linalg.inv(new_x.T @ new_x) @ new_x.T @ y  
    return beta
```

```
beta = multilinear_regression(x_train, y_train)  
beta
```

```
array([[ 8.95958970e+01],  
       [-6.77105227e+01],  
       [-2.94855410e-02],  
       [-1.25763624e-01],  
       [ 4.11175650e+00],  
       [ 2.08259777e+01]])
```

```
def prediction(x, beta):  
    b0 = np.ones(x.shape[0])  
    new_x = np.concatenate((b0.reshape(-1,1), x), axis = 1)  
    predicted_value = new_x @ beta  
    return predicted_value
```

```
predicted_y1 = prediction(x_test, beta)
```

```
ss_res = np.sum((predicted_y1 - y_test)**2)  
ss_tot = np.sum((predicted_y1 - y_test.mean())**2)  
R_sq = 1 - (ss_res/ss_tot)  
mse = np.mean((predicted_y1 - y_test)**2)  
print('R-squared Value =', R_sq)
```

R-squared Value = 0.8910828666212713

```
from sklearn.linear_model import LinearRegression  
from sklearn.metrics import mean_squared_error, r2_score  
LR = LinearRegression()  
#Training  
LR.fit(x_train, y_train)  
#Predicting  
y_pred_linear = LR.predict(x_test)
```

```
#Finding the goodness of fit  
mse_linear = mean_squared_error(y_test, y_pred_linear)  
R_sq_linear = r2_score(y_test, y_pred_linear)  
print('Mean-squared Error =', mse_linear)  
print('R-squared Value =', R_sq_linear)
```

Mean-squared Error = 10.93388128839549

R-squared Value = 0.8955413449874704

```

from sklearn.linear_model import Ridge
ridge = Ridge(alpha = 0.2)
#Training
ridge.fit(x_train, y_train)
#Predicting
y_pred_ridge = ridge.predict(x_test)

#Finding the goodness of fit
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
R_sq_ridge = r2_score(y_test, y_pred_ridge)
print('Mean-squared Error =', mse_ridge)
print('R-squared Value =', R_sq_ridge)

Mean-squared Error = 11.018535552002326
R-squared Value = 0.8947325863880121

from sklearn.linear_model import Lasso
lasso = Lasso(alpha = 0.2)
#Training
lasso.fit(x_train, y_train)
#Predicting
y_pred_lasso = lasso.predict(x_test)

#Finding the goodness of fit
mse_lasso = mean_squared_error(y_test, y_pred_lasso)
R_sq_lasso = r2_score(y_test, y_pred_lasso)
print('Mean-squared Error =', mse_lasso)
print('R-squared Value =', R_sq_lasso)

Mean-squared Error = 12.64835459163122
R-squared Value = 0.8791618388828112

data = {
    'Metrics' : ['R2 Score', 'MSE', 'RMSE'],
    'Scratch' : [R_sq, mse, (mse)**0.5],
    'sklearn' : [R_sq_linear, mse_linear, (mse_linear)**0.5],
    'Ridge' : [R_sq_ridge, mse_ridge, (mse_ridge)**0.5],
    'Lasso' : [R_sq_lasso, mse_lasso, (mse_lasso)**0.5]
}
data = pd.DataFrame(data)
data = data.set_index('Metrics')
data


```

| | Scratch | sklearn | Ridge | Lasso |
|----------|-----------|-----------|-----------|-----------|
| Metrics | | | | |
| R2 Score | 0.891083 | 0.895541 | 0.894733 | 0.879162 |
| MSE | 10.933881 | 10.933881 | 11.018536 | 11.018536 |
| RMSE | 3.306642 | 3.306642 | 3.319418 | 3.556453 |

```

plt.figure(figsize=(12, 6))

```



```

# Manual Model
plt.subplot(1,2,1)
plt.scatter(y_test, predicted_y1, alpha = 0.4)
plt.title('Manual Model')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'k--', lw=0.8)

# sklearn Model
plt.subplot(1,2,2)
plt.scatter(y_test, y_pred_linear, alpha = 0.4, c='red')
plt.title('sklearn Model')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
'k--', lw=0.8)

[<matplotlib.lines.Line2D at 0x12ae79a4c10>]

```



```

correlation = df2.corrwith(df2['Y1'])
correlation

```

```

X1    0.622272
X2   -0.658120
X3    0.455671
X4   -0.861828
X5    0.889430

```



```
X6    -0.002587
X7     0.269842
X8     0.087368
Y1     1.000000
Y2     0.975862
dtype: float64
```

Recommendations

Insulation : apparantly has a positive impact on heating load, so more the insulation less the heating demand.

Relative Compactness: Large negative coefficient ($-2.549774e+01$) indicates that more compact buildings significantly reduce heating load. Suggestion: Design compact structures to minimize exposed surface area and maximize energy efficiency.

Surface Area: Large positive coefficient ($1.139458e+15$) indicates that increasing surface area dramatically increases the heating load. Suggestion: Optimize the surface area to reduce heat loss. Avoid designs with unnecessarily large exposed areas. Use materials with better insulation properties.

X3 should be optimized for better energy performance.