

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
data = pd.read_csv('data_melbourne.csv')
```

```
data.head()
```

Unnamed: 0	Average Outflow	Average Inflow	Energy Consumption
0	2.941	2.589	175856
1	2.936	2.961	181624
2	2.928	3.225	202016
3	2.928	3.354	207547
4	2.917	3.794	202824

Biological Oxygen Demand	Chemical Oxygen Demand	Total Nitrogen
365.0	730.0	60.378
370.0	740.0	60.026
418.0	836.0	64.522
430.0	850.0	63.000
508.0	1016.0	65.590

Average Temperature	Maximum temperature	Minimum temperature
19.3	25.1	12.6
17.1	23.6	12.3
16.8	27.2	8.8
14.6	19.9	11.1
13.4	19.1	8.0

Atmospheric pressure	Average humidity	Total rainfall	Average visibility
0.0	56	1.52	
0.0	63	0.00	
0.0	47	0.25	
0.0	49	0.00	
0.0	65	0.00	

Average wind speed	Maximum wind speed	Year	Month	Day
--------------------	--------------------	------	-------	-----

0	26.9	53.5	2014	1	1
1	14.4	27.8	2014	1	2
2	31.9	61.1	2014	1	5
3	27.0	38.9	2014	1	6
4	20.6	35.2	2014	1	7

#We have dropped the first column

```
data = pd.DataFrame(data)
```

```
data.head()
```

	0	1	2	3	4	5	6
0	0.371338	0.000000	0.210224	0.1750	0.316901	0.276119	0.391885
1	0.370707	0.022712	0.230700	0.1500	0.323944	0.283582	0.385115
2	0.369697	0.038830	0.303092	0.3625	0.391549	0.355224	0.471577
3	0.369697	0.046706	0.322727	0.2875	0.408451	0.365672	0.442308
4	0.368308	0.073570	0.305960	0.4125	0.518310	0.489552	0.492115

	7	8	9	10	11	12	13
0	0.543662	0.577011	0.478689	0.0	0.577320	0.084304	0.019531
1	0.481690	0.542529	0.468852	0.0	0.649485	0.000000	0.019531
2	0.473239	0.625287	0.354098	0.0	0.484536	0.013866	0.019531
3	0.411268	0.457471	0.429508	0.0	0.505155	0.000000	0.019531
4	0.377465	0.439080	0.327869	0.0	0.670103	0.000000	0.019531

	15	16	17	18
0	0.640719	0.0	0.0	0.000000
1	0.332934	0.0	0.0	0.033333
2	0.731737	0.0	0.0	0.133333
3	0.465868	0.0	0.0	0.166667
4	0.421557	0.0	0.0	0.200000

```
data.shape
```


```
(1382, 19)
```

```
data.isnull().sum()
```

```
0      0
1      0
2      0
3      0
4      0
5      0
6      0
7      0
8      0
9      0
10     0
11     0
12     0
13     0
14     0
15     0
16     0
17     0
18     0
dtype: int64

#No null columns

corr_mat = data.corr()
plt.figure(figsize=(20,20))
sns.heatmap(corr_mat, annot = True, linewidth = 1, linecolor =
'black')
plt.show()
```




```
data['Optimality'] = 10
low = data['Energy Consumption'].quantile(0.05)
high = data['Energy Consumption'].quantile(0.55)
data.loc[(data['Energy Consumption']>=low) & (data['Energy Consumption']<=high), 'Optimality'] = 1
data.loc[~((data['Energy Consumption']>=low) & (data['Energy Consumption']<=high)), 'Optimality'] = 0

data.head()
```

```
print(low)
print(high)
```

```
198111.05
281782.25
```

```
plt.boxplot(data)
plt.show()
```



```
from sklearn.model_selection import train_test_split

#Selecting Features
x = data.drop([ 'Year', 'Month', 'Day', 'Optimality', ], axis=1)
y = data['Optimality']

x_train, x_test, y_train, y_test = train_test_split(x, y,
random_state=42, test_size=0.2)

from sklearn.preprocessing import MinMaxScaler
sc = MinMaxScaler()
x_train_sc = sc.fit_transform(x_train)
x_test_sc = sc.fit_transform(x_test)

from sklearn.metrics import f1_score
```

```

#Sigmoid Function
def sigmoid(z):
    return 1/(1+np.exp(-z))

#Cost Function (Based on log-loss)
def Cost_Function(x, y, beta):
    m = len(y)
    h = sigmoid(x @ beta)
    J = -1/m * np.sum(y * np.log(h) + (1-y) * np.log(1-h))
    return J

#Gradient Descent
def gradient_descent(x, y, beta, alpha, iters):
    m = len(y)
    for i in range(iters):
        h = sigmoid(np.dot(x, beta))
        beta = beta - alpha * 1/m * (x.T @ (h - y))
    return beta

#Prediction Function
def predict(x, beta):
    probability = sigmoid(x @ beta)
    return [1 if prob >= 0.5 else 0 for prob in probability]

#Logistic Regression
def logistic_regression(x_train_sc, x_test_sc, y_train, y_test,
alpha=0.1, iters=1000):

    x_train_sc = np.c_[np.ones((x_train_sc.shape[0], 1)), x_train_sc]
    x_test_sc = np.c_[np.ones((x_test_sc.shape[0], 1)), x_test_sc]

    print("Shape of x_train_scaled:", x_train_sc.shape)
    print("Shape of x_test_scaled:", x_test_sc.shape)

    # Initialize beta to zero
    beta = np.zeros(x_train_sc.shape[1])
    print("Shape of beta:", beta.shape)

    # Train the model
    weight = gradient_descent(x_train_sc, y_train, beta, alpha, iters)

    # Make predictions on the test set
    predictions = predict(x_test_sc, weight)

    return weight, predictions

weight, predictions = logistic_regression(x_train_sc, x_test_sc,
y_train, y_test)
print("\nWeights:", weight)

```

```
# Model Evaluation
```

```
accuracy_scratch = np.mean(predictions == y_test)
```

```
accuracy_scratch
```

```
F1_scratch = f1_score(y_test, predictions)
```

```
F1_scratch
```

```
Shape of x_train_scaled: (1105, 18)
```

```
Shape of x_test_scaled: (277, 18)
```

```
Shape of beta: (18,)
```

```
Weights: [ 0.4453377 -0.6165742 -0.23130351 -0.13757097 -2.93782944
```

```
0.33700166
```

```
0.58019296 0.22046395 0.39704644 0.65456609 0.83166099
```

```
0.52282195
```

```
0.01062281 0.02225452 0.07821289 0.06448857 0.10463371 0.161629
```

```
]
```

```
0.7846153846153846
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score
```

```
logr= LogisticRegression()
```

```
logr.fit(x_train_sc, y_train)
```

```
y_pred_logr = logr.predict(x_test_sc)
```

```
accuracy_logr = accuracy_score(y_test, y_pred_logr)
```

```
accuracy_logr
```

```
F1_logr = f1_score(y_test, y_pred_logr)
```

```
F1_logr
```

```
0.8384615384615385
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
knn= KNeighborsClassifier(n_neighbors = 10)
```

```
knn.fit(x_train_sc, y_train)
```

```
y_pred_knn = knn.predict(x_test_sc)
```

```
accuracy_knn = accuracy_score(y_test, y_pred_knn)
```

```
accuracy_knn
```

```
F1_knn = f1_score(y_test, y_pred_knn)
```

```
F1_knn
```

```
0.8045977011494253
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
dtc = DecisionTreeClassifier(random_state=42)
```

```

dtc.fit(x_train_sc, y_train)

y_predicted = dtc.predict(x_test_sc)

accuracy_dtc = accuracy_score(y_test, y_predicted)
accuracy_dtc
F1_dtc = f1_score(y_test, y_predicted)
F1_dtc

0.8957528957528957

from sklearn.ensemble import RandomForestClassifier
rf_classifier = RandomForestClassifier(random_state=42)

#Training the model
rf_classifier.fit(x_train_sc, y_train)

#Predictions
y_pred_rf = rf_classifier.predict(x_test_sc)

#Model Evaluation
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print("Random Forest Classifier- \nAccuracy:", accuracy_rf)
F1_rf = f1_score(y_test, y_pred_rf)
print(f"F1 Score: {F1_rf:.2f}")

Random Forest Classifier-
Accuracy: 0.8989169675090253
F1 Score: 0.89

from sklearn.svm import SVC
svc_classifier = SVC(kernel='rbf')

#Training the model
svc_classifier.fit(x_train_sc, y_train)

# Make predictions
y_pred_svc = svc_classifier.predict(x_test_sc)

#Model Evaluation
accuracy_svc = accuracy_score(y_test, y_pred_svc)
print("Support Vector Classifier- \nAccuracy:", accuracy_svc)
F1_svc = f1_score(y_test, y_pred_svc)
print(f"F1 Score: {F1_svc:.2f}")

Support Vector Classifier-
Accuracy: 0.8989169675090253
F1 Score: 0.90

from sklearn.metrics import confusion_matrix
confusion_scratch = confusion_matrix(y_test, predictions)

```



```

confusion_sklearn = confusion_matrix(y_test, y_pred_logr)
confusion_knn = confusion_matrix(y_test, y_pred_knn)
confusion_dt = confusion_matrix(y_test, y_pred_dtc)
confusion_rf = confusion_matrix(y_test, y_pred_rf)
confusion_svm = confusion_matrix(y_test, y_pred_svc)
df = {
    'Metrics' : ['Accuracy', 'F1-Score', 'Confusion Matrix'],
    'LR Scratch' : [accuracy_scratch, F1_scratch, confusion_scratch],
    'LR sklearn' : [accuracy_sklearn, F1_sklearn, confusion_sklearn],
    'KNN' : [accuracy_knn, F1_knn, confusion_knn],
    'Decision Tree' : [accuracy_dtc, F1_dtc, confusion_dt],
    'Random Forest' : [accuracy_rf, F1_rf, confusion_rf],
    'SVM' : [accuracy_svc, F1_svc, confusion_svm]
}
df = pd.DataFrame(df)
df = df.set_index('Metrics')
df

```

	LR Scratch	LR sklearn \
Metrics		
Accuracy	0.797834	0.848375
F1-Score	0.784615	0.838462
Confusion Matrix	[[119, 18], [38, 102]]	[[126, 11], [31, 109]]

	KNN	Decision Tree \
Metrics		
Accuracy	0.815884	0.902527
F1-Score	0.804598	0.895753
Confusion Matrix	[[121, 16], [35, 105]]	[[134, 3], [24, 116]]

	Random Forest	SVM
Metrics		
Accuracy	0.898917	0.898917
F1-Score	0.892308	0.9
Confusion Matrix	[[133, 4], [24, 116]]	[[123, 14], [14, 126]]

```

from sklearn.model_selection import GridSearchCV
param_grid_dt = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [None, 3, 5, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': [None, 'sqrt', 'log2']
}

grid_search_dt =
GridSearchCV(estimator=DecisionTreeClassifier(random_state=42),
param_grid=param_grid_dt, cv=5, scoring='accuracy', n_jobs=-1,
verbose=1)
grid_search_dt.fit(x_train_sc, y_train)

```

```
best_params_dt = grid_search_dt.best_params_  
print("Best Hyperparameters for Decision Tree:", best_params_dt)
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits
Best Hyperparameters for Decision Tree: {'criterion': 'gini',
'max_depth': None, 'max_features': None, 'min_samples_leaf': 1,
'min_samples_split': 2}

#KNN

```
param_grid_knn = {  
    'n_neighbors': [3, 5, 7, 9, 11],  
    'weights': ['uniform', 'distance'],  
    'metric': ['euclidean', 'manhattan', 'minkowski']  
}
```

```
grid_search_knn = GridSearchCV(estimator=KNeighborsClassifier(),  
param_grid=param_grid_knn, cv=5, scoring='accuracy', n_jobs=-1,  
verbose=1)
```

```
grid_search_knn.fit(x_train_sc, y_train)
```

```
best_params_knn = grid_search_knn.best_params_
```

```
print("Best Hyperparameters for KNN:", best_params_knn)
```

Fitting 5 folds for each of 30 candidates, totalling 150 fits
Best Hyperparameters for KNN: {'metric': 'manhattan', 'n_neighbors':
11, 'weights': 'uniform'}

```
param_grid_dt = {  
    'criterion': ['gini', 'entropy'],  
    'max_depth': [None, 3, 5, 10],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4],  
    'max_features': [None, 'sqrt', 'log2']  
}
```

```
grid_search_dt =
```

```
GridSearchCV(estimator=DecisionTreeClassifier(random_state=42),  
param_grid=param_grid_dt, cv=5, scoring='accuracy', n_jobs=-1,  
verbose=1)
```

```
grid_search_dt.fit(x_train_sc, y_train)
```

```
best_params_dt = grid_search_dt.best_params_
```

```
print("Best Hyperparameters for Decision Tree:", best_params_dt)
```

Fitting 5 folds for each of 216 candidates, totalling 1080 fits
Best Hyperparameters for Decision Tree: {'criterion': 'gini',
'max_depth': None, 'max_features': None, 'min_samples_leaf': 1,
'min_samples_split': 2}

```
param_grid_rf = {  
    'n_estimators': [50, 100, 200],  
    'max_depth': [None, 10, 20, 30],  
    'min_samples_split': [2, 5, 10],
```

```

    'min_samples_leaf': [1, 2, 4],
    'max_features': ['sqrt', 'log2', None],
    'bootstrap': [True, False]
}

grid_search_rf =
GridSearchCV(estimator=RandomForestClassifier(random_state=42),
param_grid=param_grid_rf, cv=5, scoring='accuracy', n_jobs=-1,
verbose=1)
grid_search_rf.fit(x_train_sc, y_train)
best_params_rf = grid_search_rf.best_params_
print("Best Hyperparameters for Random Forest:", best_params_rf)

Fitting 5 folds for each of 648 candidates, totalling 3240 fits

#SVC
param_grid_svc = {
    'C': [0.1, 1, 10, 100],
    'kernel': ['linear', 'rbf', 'poly'],
    'gamma': ['scale', 'auto', 0.01, 0.1, 1],
    'degree': [2, 3, 4],
    'class_weight': [None, 'balanced']
}

grid_search_svc = GridSearchCV(estimator=SVC(random_state=42),
param_grid=param_grid_svc, cv=5, scoring='accuracy', n_jobs=-1,
verbose=1)
grid_search_svc.fit(x_train_scaled, y_train)
best_params_svc = grid_search_svc.best_params_
print("Best Hyperparameters for SVC:", best_params_svc)

```